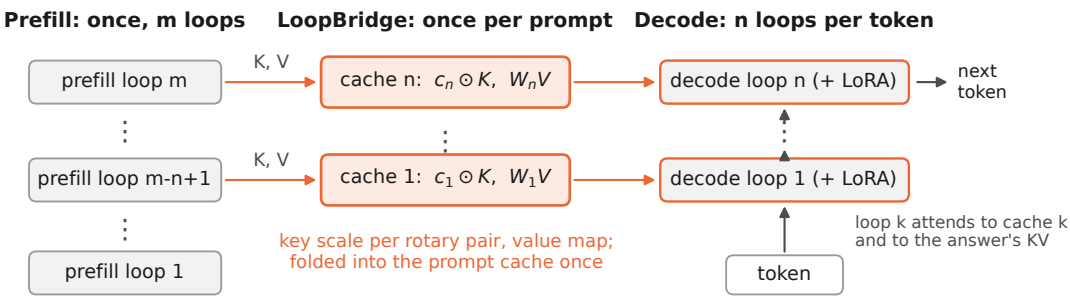
Prefill: once, m loops
LoopBridge: once per prompt
Decode: n loops per token
prefill loop m
prefill loop m-n+1
prefill loop 1
K, V
K, V
cache n: $c_n \odot K,\ W_nV$
cache 1: $c_1 \odot K,\ W_1V$
key scale per rotary pair, value map;
folded into the prompt cache once
decode loop n (+ LoRA)
decode loop 1 (+ LoRA)
next
token
loop k attends to cache k
and to the answer's KV
token