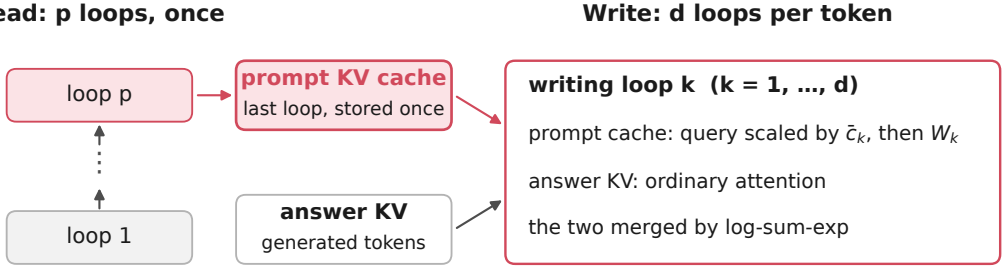
Read: p loops, once
Write: d loops per token
loop p
loop 1
prompt KV cache
last loop, stored once
answer KV
generated tokens
writing loop k (k = 1, …, d)
prompt cache: query scaled by $\tilde{c}_k$, then W_k
answer KV: ordinary attention
the two merged by log-sum-exp