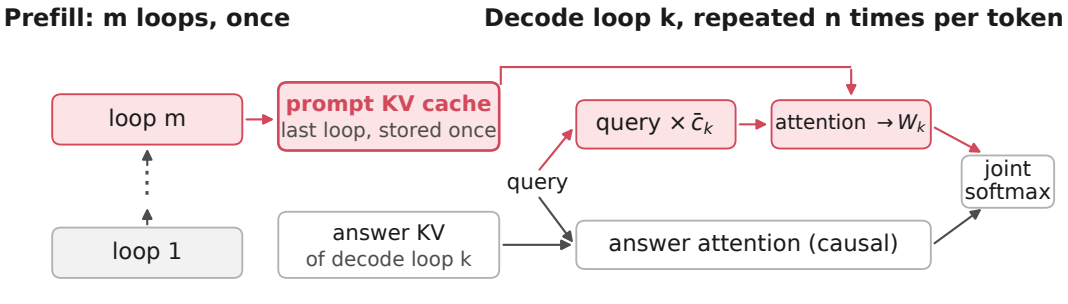
Prefill: m loops, once
Decode loop k, repeated n times per token
loop m
prompt KV cache
last loop, stored once
query × $\bar{c}_k$
attention → W_k
joint softmax
query
loop 1
answer KV
of decode loop k
answer attention (causal)