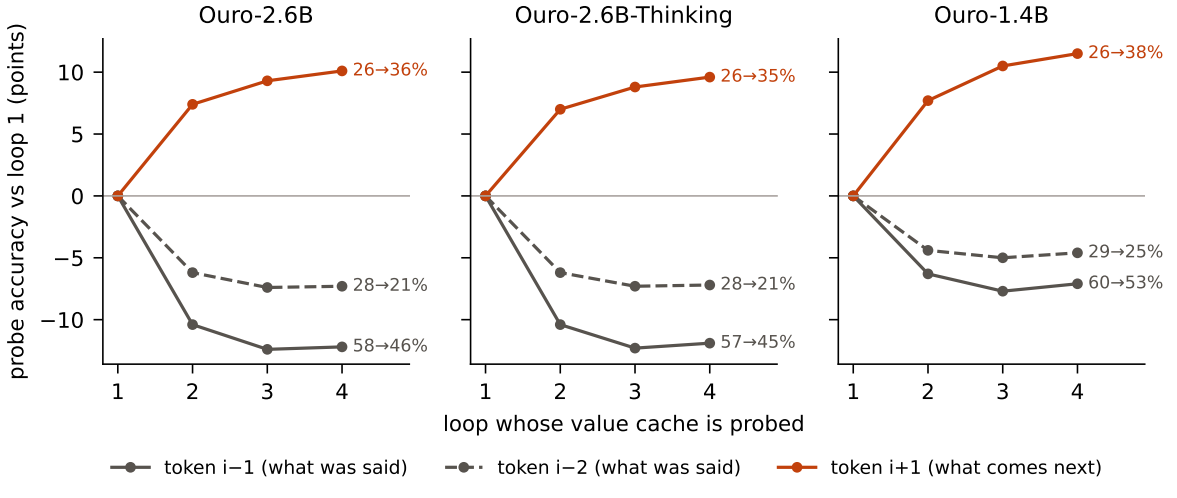
Linear probes on the stored value cache (same weights, same layer): the literal tokens fade, the next token emerges
Ouro-2.6B
Ouro-2.6B-Thinking
Ouro-1.4B
probe accuracy vs loop 1 (points)
10
5
0
−5
−10
26→36%
28→21%
58→46%
26→35%
28→21%
57→45%
26→38%
29→25%
60→53%
1
2
3
4
loop whose value cache is probed
token i−1 (what was said)
token i−2 (what was said)
token i+1 (what comes next)